

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [3]: data.describe()
```

```
Out[3]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [4]: data.head()
```

```
Out[4]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>0</b>	1	lounge	51	882	25000	1	44.907242	8.611560	8900
<b>1</b>	2	pop	51	1186	32500	1	45.666359	12.241890	8800
<b>2</b>	3	sport	74	4658	142228	1	45.503300	11.417840	4200
<b>3</b>	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
<b>4</b>	5	pop	73	3074	106880	1	41.903221	12.495650	5700

```
In [5]: data1=data.drop(['lat','lon','ID'],axis=1)
data1
```

```
Out[5]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [6]: data2=pd.get_dummies(data1)
data2
```

```
Out[6]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [7]: data2.shape
```

```
Out[7]: (1538, 8)
```

```
In [8]: y=data2['price']
x=data2.drop('price',axis=1)
```

In [9]:

y

Out[9]:

0	8900
1	8800
2	4200
3	6000
4	5700
	...
1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1538, dtype: int64

In [10]: `from sklearn.model_selection import train_test_split`  
`x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)`

In [11]: `x_test.head(5)`

Out[11]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

In [12]: `y_test.head(5)`

Out[12]:

481	7900
76	7900
1502	9400
669	8500
1409	9700

Name: price, dtype: int64

```
In [13]: from sklearn.linear_model import LinearRegression
reg=LinearRegression() #creating object of LinearRegression
reg.fit(x_train,y_train)#training and fitting LR object using training data
```

Out[13]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [14]: ypred=reg.predict(x_test)
ypred
```

Out[14]: array([ 5867.6503378 , 7133.70142341, 9866.35776216, 9723.28874535,  
10039.59101162, 9654.07582608, 9673.14563045, 10118.70728123,  
9903.85952664, 9351.55828437, 10434.34963575, 7732.26255693,  
7698.67240131, 6565.95240435, 9662.90103518, 10373.20344286,  
9599.94844451, 7699.34400418, 4941.33017994, 10455.2719478 ,  
10370.51555682, 10391.60424404, 7529.06622456, 9952.37340054,  
7006.13845729, 9000.1780961 , 4798.36770637, 6953.10376491,  
7810.39767825, 9623.80497535, 7333.52158317, 5229.18705519,  
5398.21541073, 5157.65652129, 8948.63632836, 5666.62365159,  
9822.1231461 , 8258.46551788, 6279.2040404 , 8457.38443276,  
9773.86444066, 6767.04074749, 9182.99904787, 10210.05195479,  
8694.90545226, 10328.43369248, 9069.05761443, 8866.7826029 ,  
7058.39787506, 9073.33877162, 9412.68162121, 10293.69451263,  
10072.49011135, 6748.5794244 , 9785.95841801, 9354.09969973,  
9507.9444386 , 10443.01608254, 9795.31884316, 7197.84932877,  
10108.31707235, 7009.6597206 , 9853.90699412, 7146.87414965,  
6417.69133992, 9996.97382441, 9781.18795953, 8515.83255277,  
8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,  
8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,  
8820.78555100, 10025.82571520, 7270.77108022, 8411.45804006])

```
In [15]: from sklearn.metrics import r2_score
r2_score(y_test,ypred)
```

Out[15]: 0.8415526986865394

```
In [16]: from sklearn.metrics import mean_squared_error #calculating MSE  
mean_squared_error(ypred,y_test)
```

Out[16]: 581887.727391353


```
In [17]: Results=pd.DataFrame(columns=['Price','Predicted'])  
Results['Price']=y_test  
Results['Predicted']=ypred  
Results=Results.reset_index()  
Results['ID']=Results.index  
Results.head(10)
```

Out[17]:

	index	Price	Predicted	ID
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
5	1414	9900	9654.075826	5
6	1089	9900	9673.145630	6
7	1507	9950	10118.707281	7
8	970	10700	9903.859527	8
9	1198	8999	9351.558284	9

```
In [20]: import matplotlib.pyplot as plt
sns.lineplot(x='ID',y='Price',data=Results.head(50))
sns.lineplot(x='ID',y='Predicted',data=Results.head(50))
plt.plot()
```

Out[20]: []



```
In [21]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
alpha=[1e-15,1e-10,1e-8,1e-4,1e-3,1e-2,1,5,10,20,30]
ridge=Ridge()
parameters={'alpha':alpha}
ridge_regressor=GridSearchCV(ridge,parameters)
ridge_regressor.fit(x_train,y_train)
```

```
Out[21]: GridSearchCV(estimator=Ridge(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20, 30]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [22]: ridge_regressor.best_params_
```

```
Out[22]: {'alpha': 30}
```

```
In [23]: ridge=Ridge(alpha=30)
ridge.fit(x_train,y_train)
ypred_ridge=ridge.predict(x_test)
```

```
In [24]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(ypred_ridge,y_test)
Ridge_Error
```

```
Out[24]: 579521.7970897449
```

```
In [25]: from sklearn.metrics import r2_score
r2_score(y_test,ypred_ridge)
```

```
Out[25]: 0.8421969385523054
```




```
In [26]: Results=pd.DataFrame(columns=['Actual','Predicted'])
Results['Actual']=y_test
Results['Predicted']=ypred_ridge
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```

```
Out[26]:
```

	index	Actual	Predicted	ID
0	481	7900	5869.741155	0
1	76	7900	7149.563327	1
2	1502	9400	9862.785355	2
3	669	8500	9719.283532	3
4	1409	9700	10035.895686	4
5	1414	9900	9650.311090	5
6	1089	9900	9669.183317	6
7	1507	9950	10115.128380	7
8	970	10700	9900.241944	8
9	1198	8999	9347.080772	9

```
In [27]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='ID',y='Actual',data=Results.head(50))
sns.lineplot(x='ID',y='Predicted',data=Results.head(50))
plt.plot()
```

Out[27]: []



```
In [28]: from sklearn.linear_model import ElasticNet
from sklearn.model_selection import GridSearchCV
elastic=ElasticNet()
parameters={'alpha':[1e-15,1e-10,1e-8,1e-4,1e-3,1e-2,1,5,10,20]}
elastic_regressor=GridSearchCV(elastic,parameters)
elastic_regressor.fit(x_train,y_train)
```

```
Out[28]: GridSearchCV(estimator=ElasticNet(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [29]: elastic_regressor.best_params_
```

```
Out[29]: {'alpha': 0.01}
```

```
In [30]: elastic=ElasticNet(alpha=.01)
elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
In [33]: elastic=ElasticNet(alpha=.01)
elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
Out[33]: 581390.7642825295
```

```
In [34]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_elastic)
```

```
Out[34]: 0.841688021120299
```


```
In [35]: Results=pd.DataFrame(columns=['Actual','Predicted'])
Results['Actual']=y_test
Results['Predicted']=ypred_ridge
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```

```
Out[35]:
```

	index	Actual	Predicted	ID
0	481	7900	5869.741155	0
1	76	7900	7149.563327	1
2	1502	9400	9862.785355	2
3	669	8500	9719.283532	3
4	1409	9700	10035.895686	4
5	1414	9900	9650.311090	5
6	1089	9900	9669.183317	6
7	1507	9950	10115.128380	7
8	970	10700	9900.241944	8
9	1198	8999	9347.080772	9

```
In [36]: import matplotlib.pyplot as plt
sns.lineplot(x='ID',y='Actual',data=Results.head(50))
sns.lineplot(x='ID',y='Predicted',data=Results.head(50))
plt.plot()
```

Out[36]: []



In [ ]: